

## Is AI suitable for forecasting Stock prices?

This is the code I implemented for the **AI for Finance** course final project. The idea is to predict the average price of the following day starting from open, close, high, low, volume, count, sentiment\_value of a day

### Libraries

A part from the usual python libraries, for this project there are 3 main types of libs we need to use:

- **Eikon Libraries** to exchange information with the Refinitiv database
- **Statsmodel API** to use OLS as the required econometric method
- **Neural Network libs**, divided in Natural Language in order to perform a Sentiment Analysis on news headlines and Neural Network in order to build our NN for forecasting price

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from math import floor, ceil
from IPython.display import clear_output #to write download progress in console
from os.path import exists

# Importing Eikon libraries for stock data
import eikon as ek
import cufflinks as cf
import configparser as cp
ek.set_app_key("8d0e9a8a1665482792d06cd4a41a00f3517de11b")

### Econometric Methods libraries, for OLS ###
import statsmodels.api as sm

### Neural Network libraries ##
# Natural Language Processing libraries
from nltk.sentiment import SentimentIntensityAnalyzer
# Recurrent Neural Network libraries
import tensorflow as tf
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR) # remove warnings in the training
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
```

In [3]:

```
def sign(num, threshold=0.1):
    '''Takes a number and performs the sign function (with a threshold)'''
    if (num <= threshold and num >= -1*threshold):
        return 0
    if (num < -1 * threshold):
        return -1
    else:
        return 1

def daily_sentiment_value(day, ticker):
    '''
    Takes a day in the form of 'YYYY-MM-DD' and gets the news headlines for that day
    It returns the sum of the sentiment of each headline
    '''
    # some days may return exeptions, so we ignore them assining a 0 sentiment value
```

```

try:
    titles = ek.get_news_headlines(f'R:{ticker} and LANGUAGE:EN',
                                   date_from=f'{day}T00:00:02',
                                   date_to=f'{day}T23:59:58', count = 100) ['text']
except:
    print(f'exeption at day {day}')
    return 0
# initialize the sentiment analyzer
sia = SentimentIntensityAnalyzer()
# compute the total sentiment value for the day
total_sent = 0
for t in titles:
    # we get 4 values for each headline (+, -, neutral, compound)
    total_sent+=sign(sia.polarity_scores(t) ["compound"])

return total_sent

```

## Get Stock Data

Since Eikon APIs allow to request a maximum of 100 headlines for each HTTP request, we need to make a request for each day in our dataset, resulting in a very slow data retreiving process (~2mins in total). (Un)fortunately, Eikon provides us headlines only from 1st April 2021, so we have a bit more than a year of data for this project

In [4]:

```

def get_data(DAYS=1,
             ticker = 'AAPL.O',
             sentiment = False,
             fromFile = False,
             train_size = 0.7,
             starting_day = "2021-04-01",
             ending_day = "2022-06-30"
             ):
    '''
    DAYS: number of days we want as lag
    ticker: stock ticker
    sentiment: if we want to include the sentiment column in the data
    It returns a dataset with the following columns:
    - date: as index
    - avg: (open+close)/2
    - avg_tomorrow: avg of the next day (for training purposes)
    - close: closing price
    - open: opening price
    - high: highest price
    - low: lowest price
    - volume: volume
    - count: count
    - sentiment: sentiment value for the day
    '''
    if (fromFile and exists(f'stock_data/{ticker}.pkl')):
        # print('data loaded from file')
        dataDF = pd.read_pickle(f'stock_data/{ticker}.pkl')
    else:
        # get the data from Eikon, that provides us headlines only from 1st April 2021
        dataDF = ek.get_timeseries(f'{ticker}',
                                   fields='*', # all fields
                                   start_date=starting_day,
                                   end_date=ending_day)

        # if the flag sentiment is True, we add the sentiment column to the data (slow process)

        if sentiment:
            d2 = dataDF
            d2.reset_index(inplace=True)
            # we get the list of dates
            d2['Date'] = d2['Date'].apply(lambda x: str(x)[:10])

            i_total = len(d2['Date'])
            sentiment_values = []
            for i, day in enumerate(d2['Date']):
                clear_output(wait=True)
                print(f'Downloading sentiment data: {i}/{i_total} [{round(i*100/i_total, 2

```

```

)})'])

    # we get the sentiment value for the day
    sentiment_values.append(daily_sentiment_value(day, ticker))
    clear_output(wait=True)

    dataDF['SENT'] = sentiment_values

    # add the column 'AVG' using (dataDF['close']+dataDF['open'])/2 to the dataDF
    dataDF['AVG'] = (dataDF['CLOSE']+dataDF['OPEN'])/2
    # shift the AVG column by DAYS to get the following day's average (for training)
    dataDF['AVG_TOMORROW'] = dataDF['AVG'].shift(-DAYS)
    # remove last row from dataDF (since we shifted the last row by DAYS)
    dataDF = dataDF[:-DAYS]
    dataDF.to_pickle(f"stock_data/{ticker}.pkl")
fig, plot = plt.subplots(figsize=(18, 9))
fig.suptitle(f'{ticker} stocks AVG {starting_day}-{ending_day}')
plot.plot(dataDF['AVG'], color='#219653', label=f'AVG {ticker} (USD)')
plot.axvline(int(len(dataDF)*train_size), color = '#A7E9A1', label = 'train/test')
plt.legend()
plt.show()

return dataDF

```

In [5]:

```

def split_data(x, y, train_size=0.7):
    """
    Split the data into training and testing sets divided according to the train_size parameter
    returns x_train, y_train, x_test, y_test
    """
    train_size = int(len(x)*train_size)

    x_train = x[:train_size]
    y_train = y[:train_size]

    x_test = x[train_size:]
    y_test = y[train_size:]

    return x_train, y_train, x_test, y_test

```

## Ordinary Least Squares

Chosen Econometric Method.

In [6]:

```

def do_OLS(x, y, x_names=[], y_name=''):
    """
    x: array of training x values
    y: array of training y values
    x_names: list of features names
    y_name: name of the target variable
    It returns the OLS model parameters
    """

    X = sm.add_constant(np.array(x, dtype='float32' ))
    model = sm.OLS(np.array(y, dtype='float32' ), X)
    fitted = model.fit()
    # print the results

    if len(x_names) == 0 or len(y_name) == 0:
        print(fitted.summary())
    else:
        print(fitted.summary2(xname=['const']+x_names, yname=y_name))

    return fitted.params

```

## Recurrent Neural Network

Manually build with 4 LSTM (Long Short-Term Memory) layers with 64 units, ending in a Dense layer with 1 unit to get the value.

In [7]:

```
def do_RNN(x_train, x_labels, y_train, x_test, y_test, epochs=10, n_steps=1):  
    """  
    Perform a RNN model training on the data, returns the list of predictions on x_test  
    """  
    # since we are working with data of different magnitudes, we need to scale and reshape  
    the data  
    scalerX = MinMaxScaler(feature_range=(0, 1))  
    scalerY = MinMaxScaler(feature_range=(0, 1))  
    x_train = scalerX.fit_transform(x_train)  
    x_test = scalerX.transform(x_test)  
    y_train = scalerY.fit_transform(y_train.reshape(-1, 1))  
  
    # we need to reshape the data to be compatible with the RNN  
    n_records_train = len(x_train)  
    n_features = len(x_labels)  
    X_train = []  
    Y_train = []  
    for i in range(n_steps, n_records_train):  
        X_train.append(x_train[i-n_steps:i])  
        Y_train.append(y_train[i][0])  
    X_train, Y_train = np.array(X_train), np.array(Y_train)  
    X_train_shaped = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], X_train.shap  
e[2]))  
  
    # we start building the model  
    model = Sequential()  
    # LSTM layers are the one with some 'memory of the past' values (Long Short Term Memory  
) , they use past values for their predictions  
    model.add(LSTM(units = 64, return_sequences = True, input_shape = (n_steps, n_features  
)))  
    model.add(LSTM(units = 64, return_sequences = True))  
    model.add(LSTM(units = 64, return_sequences = True))  
    model.add(LSTM(units = 64, return_sequences = True))  
    # One 'neuron' for the output since we want a number  
    model.add(Dense(units = 1))  
    # we use the mean squared error as the loss function and Adam as optimizer since it all  
ows to get better results quickly and then it settles  
    model.compile(optimizer='adam', loss='mean_squared_error')  
    # we train the model  
    model.fit(X_train_shaped, Y_train, epochs=epochs, batch_size = 32, verbose = 0)  
  
    # we get the predictions on the test set  
    n_records_test = len(x_test)  
    X_test = []  
    for i in range(n_steps, n_records_test):  
        X_test.append(x_test[i-n_steps:i])  
    X_test = np.array(X_test)  
    y_res = model.predict(X_test)  
    predictions_transformed = []  
    for val in y_res:  
        predictions_transformed.append(val[0])  
    predictions_transformed = np.array(predictions_transformed)  
    # we need to inverse the scaling of the predictions and we return them  
    return scalerY.inverse_transform(predictions_transformed)
```

## MAIN Function

- Get data from Eikon (or from file if it's available)
- Perform Sentiment Analysis on news headline
- Compute OLS Regression and RNN regression
- Compare results on test
- Compute MSE

In [8]:

```

def main(
    epochs = 40,
    lag = 1,
    train_size = 0.7,
    sentiment=True,
    ticker = 'AAPL.O',
    fromFile = True
):
    # get data from Eikon database
    data = get_data(sentiment=sentiment, ticker=ticker, fromFile=fromFile, train_size=train_size)
    # if we don't want to download sentiment data, we just set sentiment=False
    if sentiment:
        x_labels = ['HIGH', 'LOW', 'OPEN', 'CLOSE', 'COUNT', 'VOLUME', 'SENT']
    else:
        x_labels = ['HIGH', 'LOW', 'OPEN', 'CLOSE', 'COUNT', 'VOLUME']
    y_label = 'AVG_TOMORROW'

    # divide and prepare data
    x = data[x_labels].values
    y = data[y_label].values
    x_train, y_train, x_test, y_test = split_data(x, y, train_size=train_size)
    # compute the OLS econometric method
    parameters = do_OLS(x_train, y_train, x_names=x_labels, y_name=y_label)
    # compute the results of the OLS model
    OLS_predictions = []
    for entry in x_test:
        tmp = 0
        for i, p in enumerate(parameters):
            if i == 0:
                tmp += p
            else:
                tmp += p*entry[i-1]
        OLS_predictions.append(tmp)
    # get results of x_test from the builded Neural Network
    res = do_RNN(x_train, x_labels, y_train, x_test, y_test, epochs=epochs, n_steps=lag)

    # prepare data to plot
    df_NN = pd.DataFrame(res, columns = ['NN']).shift(-lag)
    df_OLS = pd.DataFrame(OLS_predictions, columns = ['OLS']).shift(-lag)
    train_size = int(len(x)*train_size)
    df_Actual = pd.DataFrame(y_test, columns = ['Actual'])
    dfResults = pd.concat([df_OLS, df_Actual, df_NN], axis=1)
    dfResults = dfResults[:-2*lag]

    fig, plot = plt.subplots(figsize=(18, 9))
    plot.plot(dfResults['Actual'], color='blue', label='Actual', linewidth=2)
    plot.plot(dfResults['OLS'], color='red', label='OLS', linewidth=1)
    plot.plot(dfResults['NN'], color='green', label='NN', linewidth=1)
    plt.legend()
    plt.show()

    # calculate mse for both NN and OLS
    mse_NN = np.mean(np.square(dfResults['NN']-dfResults['Actual']))
    mse_OLS = np.mean(np.square(dfResults['OLS']-dfResults['Actual']))
    print('Mean Squared Error of NN model: ', mse_NN)
    print('Mean Squared Error of OLS: ', mse_OLS)

```

In [9]:

```
main(ticker = 'AAPL.O')
```

AAPL.O stocks AVG 2021-04-01-2022-06-30






# Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.991
Dependent Variable:   AVG_TOMORROW        AIC:              819.9010
Date:                2022-07-11 17:44    BIC:              847.0136
No. Observations:    219                Log-Likelihood:    -401.95
Df Model:             7                  F-statistic:       3291.
Df Residuals:         211                Prob (F-statistic): 1.29e-211
R-squared:            0.991              Scale:           2.3872
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	1.6959	1.1537	1.4699	0.1431	-0.5784	3.9702
HIGH	0.0695	0.1818	0.3822	0.7027	-0.2889	0.4279
LOW	0.3605	0.1777	2.0292	0.0437	0.0103	0.7107
OPEN	-0.2663	0.1421	-1.8745	0.0622	-0.5463	0.0137
CLOSE	0.8248	0.1158	7.1241	0.0000	0.5965	1.0530
COUNT	0.0000	0.0000	2.1213	0.0351	0.0000	0.0000
VOLUME	-0.0000	0.0000	-1.7591	0.0800	-0.0000	0.0000
SENT	0.0211	0.0137	1.5408	0.1249	-0.0059	0.0482

```
=====
Omnibus:              25.455              Durbin-Watson:       1.804
Prob(Omnibus):        0.000              Jarque-Bera (JB):    89.093
Skew:                 0.354              Prob(JB):            0.000
Kurtosis:             6.043              Condition No.:       997698643
=====
```


\* The condition number is large (1e+09). This might indicate strong multicollinearity or other numerical problems.



```
Mean Squared Error of NN model:  0.3078726244338866
Mean SquaredError of OLS:  2.709550415781781
```

```
In [10]:  
main(ticker = 'AMZN.O')
```

AMZN.O stocks AVG 2021-04-01-2022-06-30




#### Results: Ordinary least squares


```
=====
Model:                OLS                Adj. R-squared:    0.932
Dependent Variable:   AVG_TOMORROW        AIC:              975.6470
Date:                2022-07-11 17:44    BIC:              1002.7596
No. Observations:    219                Log-Likelihood:    -479.82
Df Model:            7                  F-statistic:      429.1
Df Residuals:        211                Prob (F-statistic): 4.90e-121
R-squared:           0.934              Scale:          4.8612
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	7.1581	3.2873	2.1775	0.0305	0.6780	13.6382
HIGH	0.2125	0.1973	1.0769	0.2828	-0.1765	0.6014
LOW	0.4856	0.1768	2.7472	0.0065	0.1372	0.8341
OPEN	-0.4611	0.1686	-2.7347	0.0068	-0.7934	-0.1287
CLOSE	0.7190	0.1502	4.7856	0.0000	0.4228	1.0151
COUNT	0.0000	0.0000	1.6785	0.0947	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-1.1119	0.2675	-0.0000	0.0000
SENT	0.0020	0.0112	0.1791	0.8581	-0.0201	0.0241

```
=====
Omnibus:              72.461            Durbin-Watson:        2.063
Prob(Omnibus):        0.000            Jarque-Bera (JB):    2043.975
Skew:                 -0.534            Prob(JB):           0.000
Kurtosis:             17.928            Condition No.:      1655746863
=====
```

\* The condition number is large (2e+09). This might indicate strong multicollinearity or other numerical problems.





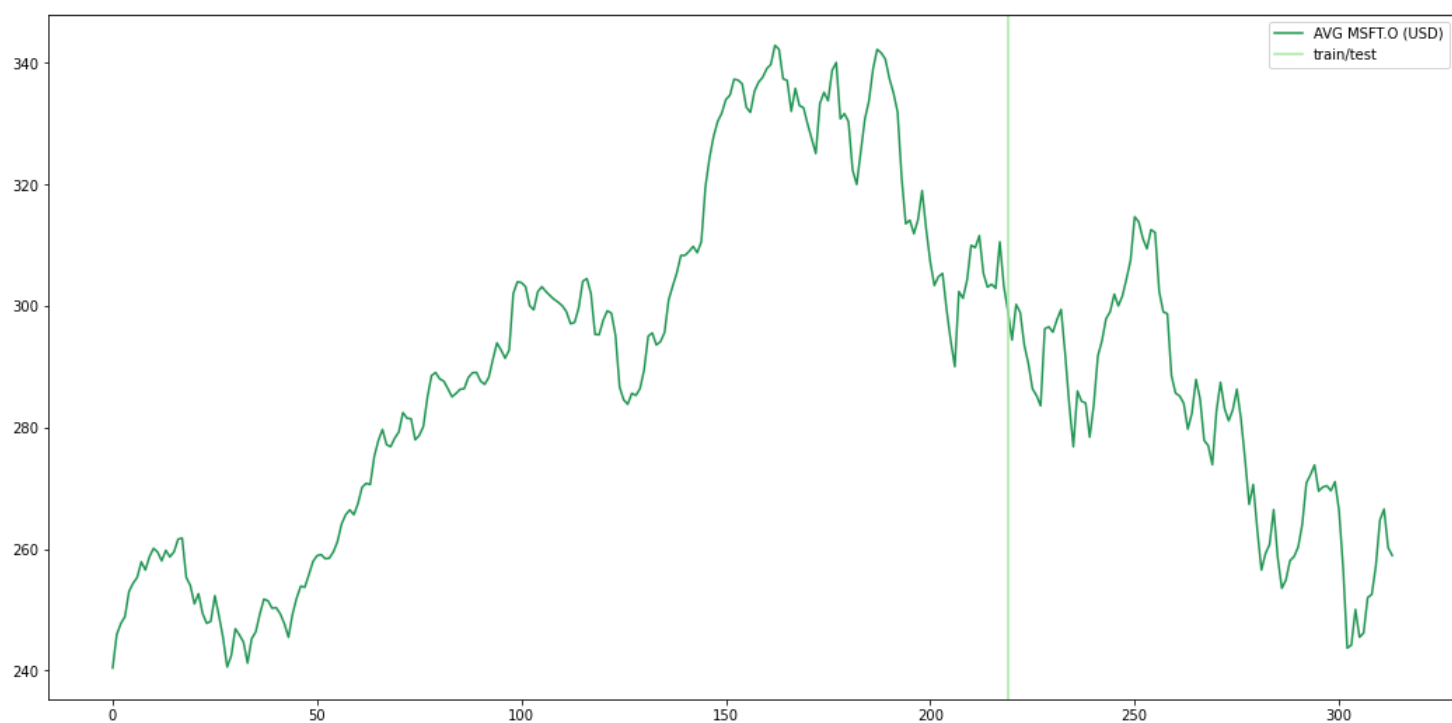
Mean Squared Error of NN model: 354.98378447301786

Mean Squared Error of OLS: 18.039098109551034

In [11]:

```
main(ticker = 'MSFT.O')
```

MSFT.O stocks AVG 2021-04-01-2022-06-30



Results: Ordinary least squares


```
=====
Model:                OLS                Adj. R-squared:      0.991
Dependent Variable:    AVG_TOMORROW        AIC:                1091.2358
Date:                 2022-07-11 17:44    BIC:                1118.3484
No. Observations:     219                Log-Likelihood:     -537.62
Df Model:              7                  F-statistic:        3284.
Df Residuals:         211                Prob (F-statistic): 1.57e-211
R-squared:             0.991              Scale:             8.2407
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	3.0106	2.2313	1.3492	0.1787	-1.3879	7.4091
HIGH	0.1915	0.1724	1.1106	0.2680	-0.1484	0.5313
LOW	0.0376	0.1609	0.2336	0.8155	-0.2795	0.3547
OPEN	-0.0648	0.1288	-0.5034	0.6152	-0.3187	0.1890
CLOSE	0.8242	0.1110	7.4224	0.0000	0.6053	1.0431
COUNT	0.0000	0.0000	0.5179	0.6051	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-0.6747	0.5006	-0.0000	0.0000
SENT	0.0510	0.0244	2.0870	0.0381	0.0028	0.0991

```
=====
Omnibus:              14.292              Durbin-Watson:       1.918
Prob(Omnibus):         0.001              Jarque-Bera (JB):    38.100
Skew:                  -0.108              Prob(JB):            0.000
Kurtosis:              5.032              Condition No.:       335860040
=====
```

\* The condition number is large (3e+08). This might indicate strong multicollinearity or other numerical problems.



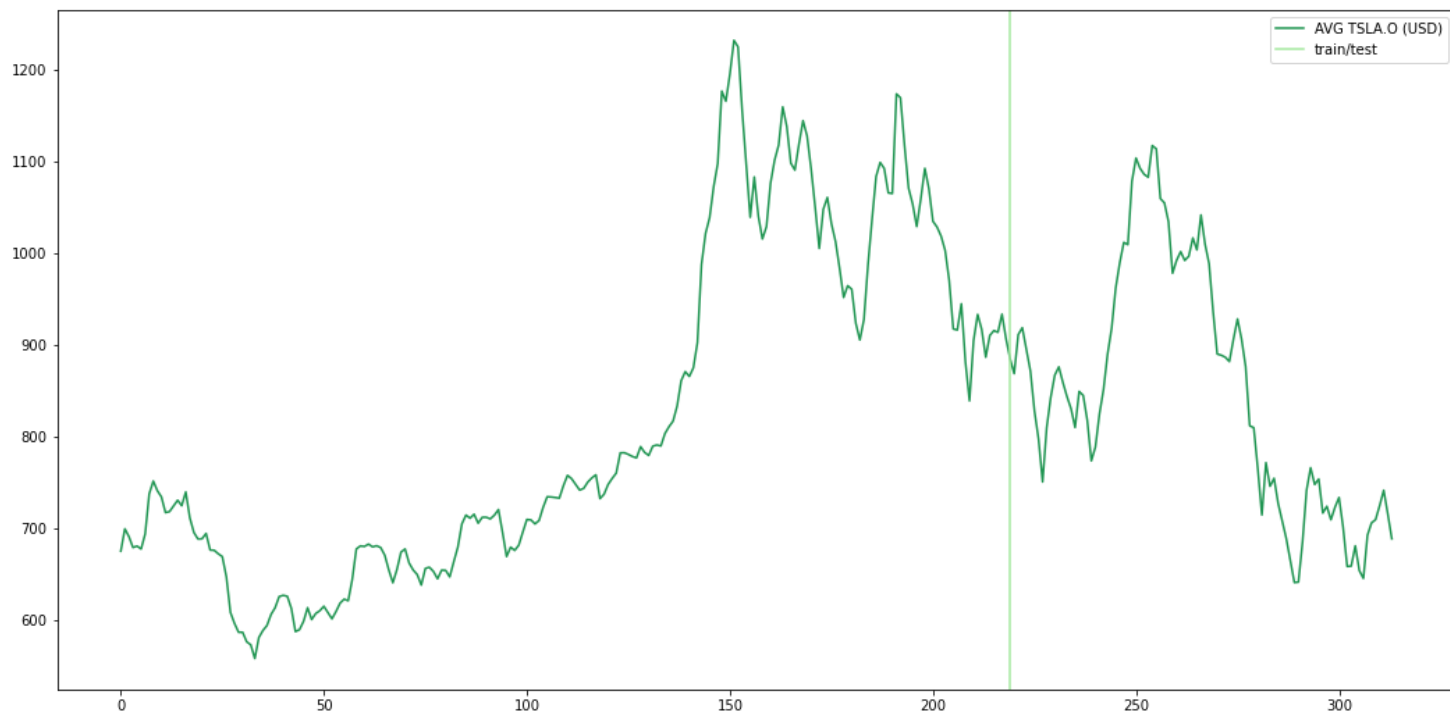


Mean Squared Error of NN model: 3.5941229033255193  
Mean Squared Error of OLS: 5.746718169131486

In [12]:

```
main(ticker = 'TSLA.O')
```

TSLA.O stocks AVG 2021-04-01-2022-06-30



Results: Ordinary least squares


```
=====
Model: OLS Adj. R-squared: 0.988
Dependent Variable: AVG_TOMORROW AIC: 1938.2710
Date: 2022-07-11 17:44 BIC: 1965.3836
No. Observations: 219 Log-Likelihood: -961.14
Df Model: 7 F-statistic: 2542.
Df Residuals: 211 Prob (F-statistic): 6.58e-200
R-squared: 0.988 Scale: 394.19
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	0.6783	10.2520	0.0662	0.9473	-19.5312	20.8877
HIGH	-0.1348	0.1760	-0.7658	0.4446	-0.4819	0.2122

LOW	0.1427	0.1619	0.8812	0.3792	-0.1765	0.4618
OPEN	-0.0057	0.1321	-0.0433	0.9655	-0.2662	0.2547
CLOSE	1.0025	0.1209	8.2913	0.0000	0.7642	1.2409
COUNT	0.0000	0.0000	0.6922	0.4896	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-0.3347	0.7382	-0.0000	0.0000
SENT	-0.4144	0.1614	-2.5677	0.0109	-0.7325	-0.0962

Omnibus:	64.807	Durbin-Watson:	1.898
Prob(Omnibus):	0.000	Jarque-Bera (JB):	352.122
Skew:	1.014	Prob(JB):	0.000
Kurtosis:	8.872	Condition No.:	206721043

\* The condition number is large (2e+08). This might indicate strong multicollinearity or other numerical problems.




Mean Squared Error of NN model: 53.66423920956459  
Mean Squared Error of OLS: 357.66524529296174

In [13]:

```
main(ticker = 'KO')
```

KO stocks AVG 2021-04-01-2022-06-30



# Results: Ordinary least squares

```

=====
Model:                OLS                Adj. R-squared:    0.979
Dependent Variable:   AVG_TOMORROW        AIC:              145.6368
Date:                2022-07-11 17:45    BIC:              172.7494
No. Observations:    219                Log-Likelihood:    -64.818
Df Model:            7                  F-statistic:       1430.
Df Residuals:        211                Prob (F-statistic): 5.62e-174
R-squared:           0.979              Scale:           0.10984
=====

```


	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	1.1257	0.5780	1.9475	0.0528	-0.0137	2.2652
HIGH	-0.0246	0.1480	-0.1661	0.8682	-0.3162	0.2671
LOW	0.0577	0.1528	0.3774	0.7062	-0.2435	0.3589
OPEN	0.0250	0.1106	0.2262	0.8213	-0.1930	0.2430
CLOSE	0.9209	0.1127	8.1698	0.0000	0.6987	1.1432
COUNT	0.0000	0.0000	1.3523	0.1777	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-1.2127	0.2266	-0.0000	0.0000
SENT	0.0001	0.0077	0.0150	0.9880	-0.0150	0.0152

```

=====
Omnibus:            31.886              Durbin-Watson:      1.923
Prob(Omnibus):      0.000              Jarque-Bera (JB):   205.973
Skew:               -0.186              Prob(JB):           0.000
Kurtosis:           7.736              Condition No.:      428713390
=====

```

\* The condition number is large (4e+08). This might indicate strong multicollinearity or other numerical problems.



```


Mean Squared Error of NN model:  0.03518033576646622
Mean Squared Error of OLS:  0.14887357320834885

```

In [14]:

```
main(ticker = 'MDT')
```

MDT stocks AVG 2021-04-01-2022-06-30






# Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.988
Dependent Variable:   AVG_TOMORROW       AIC:                653.1308
Date:                2022-07-11 17:45   BIC:                680.2433
No. Observations:    219                Log-Likelihood:     -318.57
Df Model:             7                  F-statistic:        2576.
Df Residuals:        211                Prob (F-statistic): 1.60e-200
R-squared:            0.988              Scale:             1.1147
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	1.8628	1.3473	1.3826	0.1682	-0.7931	4.5188
HIGH	0.0915	0.1373	0.6664	0.5059	-0.1791	0.3621
LOW	-0.3642	0.1725	-2.1114	0.0359	-0.7043	-0.0242
OPEN	0.0964	0.1154	0.8353	0.4045	-0.1311	0.3239
CLOSE	1.1627	0.1289	9.0172	0.0000	0.9085	1.4168
COUNT	-0.0000	0.0000	-0.8876	0.3758	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-0.5657	0.5722	-0.0000	0.0000
SENT	0.0218	0.0250	0.8704	0.3851	-0.0275	0.0711

```
=====
Omnibus:              117.048            Durbin-Watson:        2.076
Prob(Omnibus):        0.000              Jarque-Bera (JB):     930.070
Skew:                 -1.920              Prob(JB):             0.000
Kurtosis:             12.337              Condition No.:        110772817
=====
```


\* The condition number is large (1e+08). This might indicate strong multicollinearity or other numerical problems.



```
Mean Squared Error of NN model:  14.32723729168111
Mean Squared Error of OLS:  0.6154328313213744
```

```
main(ticker = 'INTC.O')
```

INTC.O stocks AVG 2021-04-01-2022-06-30




## Results: Ordinary least squares


```
=====
Model:                OLS                Adj. R-squared:    0.968
Dependent Variable:    AVG_TOMORROW        AIC:              475.7126
Date:                 2022-07-11 17:45    BIC:              502.8251
No. Observations:     219                Log-Likelihood:    -229.86
Df Model:              7                  F-statistic:       951.7
Df Residuals:         211                Prob (F-statistic): 8.25e-156
R-squared:            0.969                Scale:           0.49583
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	2.2691	0.7483	3.0324	0.0027	0.7941	3.7442
HIGH	0.0143	0.1918	0.0748	0.9405	-0.3638	0.3925
LOW	-0.1066	0.2020	-0.5277	0.5983	-0.5047	0.2916
OPEN	0.1743	0.1504	1.1590	0.2477	-0.1222	0.4709
CLOSE	0.8812	0.1454	6.0618	0.0000	0.5946	1.1677
COUNT	0.0000	0.0000	0.6925	0.4894	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-1.6250	0.1057	-0.0000	0.0000
SENT	-0.0284	0.0116	-2.4468	0.0152	-0.0512	-0.0055

```
=====
Omnibus:              173.845            Durbin-Watson:      1.860
Prob(Omnibus):        0.000              Jarque-Bera (JB):   4687.876
Skew:                 -2.708              Prob(JB):          0.000
Kurtosis:             25.009              Condition No.:     495588717
=====
```

\* The condition number is large (5e+08). This might indicate strong multicollinearity or other numerical problems.






Mean Squared Error of NN model: 9.014551307937955  
Mean Squared Error of OLS: 0.16681711345034678

In [16]:

```
main(ticker = 'AMD.O')
```

AMD.O stocks AVG 2021-04-01-2022-06-30




Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.993
Dependent Variable:    AVG_TOMORROW        AIC:                956.7436
Date:                 2022-07-11 17:45    BIC:                983.8562
No. Observations:     219                Log-Likelihood:     -470.37
Df Model:              7                  F-statistic:        4139.
Df Residuals:          211                Prob (F-statistic): 4.85e-222
R-squared:             0.993              Scale:             4.4592
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	-0.4790	0.8195	-0.5845	0.5595	-2.0945	1.1365
HIGH	-0.2950	0.1510	-1.9534	0.0521	-0.5927	0.0027
LOW	0.2781	0.1511	1.8396	0.0672	-0.0199	0.5760
OPEN	-0.0742	0.1113	-0.6668	0.5057	-0.2937	0.1452
CLOSE	1.1045	0.1058	10.4403	0.0000	0.8959	1.3130
COUNT	0.0000	0.0000	1.1296	0.2599	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-1.2376	0.2172	-0.0000	0.0000
SENT	0.1301	0.0356	3.6507	0.0003	0.0598	0.2003

```
=====
Omnibus:              16.219              Durbin-Watson:        2.164
Prob(Omnibus):         0.000              Jarque-Bera (JB):     26.148
Skew:                  0.434              Prob(JB):             0.000
Kurtosis:              4.453              Condition No.:        377766364
=====
```

\* The condition number is large (4e+08). This might indicate strong multicollinearity or other numerical problems.




Mean Squared Error of NN model: 0.6372774357276006  
Mean Squared Error of OLS: 5.339673521376355

In [17]:

```
main(ticker = 'HAS.O')
```

HAS.O stocks AVG 2021-04-01-2022-06-30



Results: Ordinary least squares


```
=====
Model: OLS Adj. R-squared: 0.923
Dependent Variable: AVG_TOMORROW AIC: 603.2424
Date: 2022-07-11 17:45 BIC: 630.3550
No. Observations: 219 Log-Likelihood: -293.62
Df Model: 7 F-statistic: 371.9
Df Residuals: 211 Prob (F-statistic): 5.86e-115
R-squared: 0.925 Scale: 0.88764
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	6.8315	1.9166	3.5645	0.0005	3.0535	10.6096
HIGH	-0.0293	0.1301	-0.2253	0.8220	-0.2858	0.2272
LOW	0.02150	0.1386	1.5519	0.1222	-0.0581	0.4882

OPEN	-0.1111	0.1083	-1.0255	0.3063	-0.3247	0.1025
CLOSE	0.8585	0.1073	7.9993	0.0000	0.6469	1.0701
COUNT	-0.0000	0.0000	-1.2890	0.1988	-0.0001	0.0000
VOLUME	0.0000	0.0000	1.7511	0.0814	-0.0000	0.0000
SENT	-0.0032	0.0227	-0.1405	0.8884	-0.0480	0.0416

Omnibus:	183.302	Durbin-Watson:	1.955
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7180.044
Skew:	2.785	Prob(JB):	0.000
Kurtosis:	30.492	Condition No.:	27345181

\* The condition number is large (3e+07). This might indicate strong multicollinearity or other numerical problems.




Mean Squared Error of NN model: 11.46175095076568  
Mean Squared Error of OLS: 1.0464642363408918

In [18]:

```
main(ticker='IBM')
```

IBM stocks AVG 2021-04-01-2022-06-30



Results: Ordinary least squares



```

=====
Model:                OLS                Adj. R-squared:    0.964
Dependent Variable:  AVG_TOMORROW        AIC:              730.6087
Date:               2022-07-11 17:45    BIC:              757.7213
No. Observations:   219                Log-Likelihood:    -357.30
Df Model:           7                  F-statistic:       841.5
Df Residuals:       211                Prob (F-statistic): 2.31e-150
R-squared:          0.965              Scale:           1.5879
=====

```


	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	4.0754	1.8123	2.2487	0.0256	0.5028	7.6480
HIGH	-0.1246	0.1595	-0.7815	0.4354	-0.4391	0.1898
LOW	-0.1064	0.1742	-0.6108	0.5420	-0.4499	0.2370
OPEN	0.1086	0.1267	0.8577	0.3921	-0.1411	0.3584
CLOSE	1.0905	0.1325	8.2329	0.0000	0.8294	1.3516
COUNT	-0.0000	0.0000	-0.3495	0.7270	-0.0000	0.0000
VOLUME	0.0000	0.0000	0.3212	0.7484	-0.0000	0.0000
SENT	0.0626	0.0218	2.8703	0.0045	0.0196	0.1056

```

=====
Omnibus:              163.619          Durbin-Watson:       1.979
Prob(Omnibus):        0.000           Jarque-Bera (JB):    3815.977
Skew:                 -2.522          Prob(JB):            0.000
Kurtosis:             22.818          Condition No.:      127697156
=====

```

\* The condition number is large (1e+08). This might indicate strong multicollinearity or other numerical problems.



```


Mean Squared Error of NN model:  0.405187505919918
Mean Squared Error of OLS:  0.846706029230757


```

In [19]:

```
main(ticker='JNJ')
```

JNJ stocks AVG 2021-04-01-2022-06-30






# Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.960
Dependent Variable:    AVG_TOMORROW        AIC:                603.9600
Date:                 2022-07-11 17:45    BIC:                631.0726
No. Observations:     219                Log-Likelihood:     -293.98
Df Model:              7                  F-statistic:        751.1
Df Residuals:          211                Prob (F-statistic): 2.40e-145
R-squared:             0.961              Scale:            0.89055
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	5.4194	2.5201	2.1505	0.0327	0.4517	10.3872
HIGH	0.0771	0.1321	0.5835	0.5602	-0.1833	0.3375
LOW	-0.1795	0.1267	-1.4167	0.1581	-0.4294	0.0703
OPEN	0.0621	0.0919	0.6757	0.5000	-0.1191	0.2433
CLOSE	1.0082	0.1006	10.0227	0.0000	0.8099	1.2065
COUNT	-0.0000	0.0000	-0.5560	0.5788	-0.0000	0.0000
VOLUME	-0.0000	0.0000	-0.1876	0.8513	-0.0000	0.0000
SENT	0.0018	0.0074	0.2408	0.8099	-0.0128	0.0164

```
=====
Omnibus:                27.566            Durbin-Watson:        2.079
Prob(Omnibus):          0.000            Jarque-Bera (JB):     58.301
Skew:                   -0.606            Prob(JB):             0.000
Kurtosis:                5.219            Condition No.:        289643708
=====
```

\* The condition number is large (3e+08). This might indicate strong multicollinearity or other numerical problems.



```
Mean Squared Error of NN model:  0.39148453583006615
Mean Squared Error of OLS:       0.8841508974201747
```